

Methodology

```
from google.colab import drive
drive.mount('/content/drive')
```

```
from google.colab import auth
auth.authenticate_user()
```

```
from google.cloud import storage
```

```
project_id = "sharp-matter-449521-u2"
!gcloud config set project {project_id}
```

Updated property [core/project].

```
!wget -P /usr/lib/spark/jars/ https://storage.googleapis.com/hadoop-lib/gcs/gcs-connector-hadoop3-latest.jar
```

--2025-04-09 22:20:49-- <https://storage.googleapis.com/hadoop-lib/gcs/gcs-connector-hadoop3-latest.jar>
Resolving storage.googleapis.com (storage.googleapis.com)... 172.253.117.207, 142.250.99.207, 142.250.107.207, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|172.253.117.207|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 40713341 (39M) [application/java-archive]
Saving to: '/usr/lib/spark/jars/gcs-connector-hadoop3-latest.jar'

gcs-connector-hadoo 100%[=====] 38.83M 207MB/s in 0.2s

2025-04-09 22:20:49 (207 MB/s) - '/usr/lib/spark/jars/gcs-connector-hadoop3-latest.jar' saved [40713341/40713341]

```
import warnings
warnings.filterwarnings('ignore')
```

```
from pyspark.sql import SparkSession
```

```
spark = SparkSession.builder \
    .appName("BigDataProcessing") \
    .config("spark.jars", "/usr/lib/spark/jars/gcs-connector-hadoop3-latest.jar") \
    .config("spark.hadoop.fs.gs.impl", "com.google.cloud.hadoop.fs.gcs.GoogleHadoopFileSystem") \
    .config("spark.hadoop.fs.gs.auth.service.account.enable", "true") \
    .getOrCreate()
```

spark

SparkSession - in-memory

SparkContext

[Spark UI](#)

Version
v3.5.5
Master
local[*]
AppName
BigDataProcessing

```
df_reviews = spark.read.parquet('gs://final_dataset_dat490/dat490_final_dataset_cleaned.parquet', headers=True, inferSchema=True)
```

```
df_reviews.columns
```

```
['gmap_id',
 'customer_name',
 'rating',
 'reviews',
 'time',
 'avg_rating',
 'category',
 'latitude',
 'longitude',
 'business_name',
 'num_of_reviews',
```

```
'state',
'standard_category',
'Monday',
'Tuesday',
'Wednesday',
'Thursday',
'Friday',
'Saturday',
'Sunday',
'timestamp',
'week',
'month',
'year',
'time_seconds']
```

✓ VADER

```
from pyspark.sql.functions import udf
from pyspark.sql.types import FloatType
from nltk.sentiment.vader import SentimentIntensityAnalyzer
import nltk
nltk.download("vader_lexicon")
```

```
# Initializing VADER
sia = SentimentIntensityAnalyzer()
```

```
def vader_sentiment(text):
    if text:
        return float(sia.polarity_scores(text)["compound"])
    else:
        return 0.0
```

```
vader_udf = udf(vader_sentiment, FloatType())
```

```
df_sentiment = df_reviews.withColumn("sentiment_score", vader_udf("reviews"))
```

↳ [nltk_data] Downloading package vader_lexicon to /root/nltk_data...

```
from pyspark.sql.functions import when
```

```
df_sentiment = df_sentiment.withColumn(
    "sentiment_label",
    when(df_sentiment["sentiment_score"] > 0.2, "Positive")
    .when(df_sentiment["sentiment_score"] < -0.2, "Negative")
    .otherwise("Neutral")
)
```

```
from pyspark.sql.functions import approx_count_distinct
```

```
df_sentiment.groupBy("sentiment_label").agg(approx_count_distinct("gmap_id").alias("approx_count")).show()
```

↳

sentiment_label	approx_count
Positive	2598785
Neutral	149531
Negative	197929

```
from pyspark.sql.functions import when
```

```
df_sentiment = df_sentiment.withColumn(
    "sentiment_label",
    when(df_sentiment["sentiment_score"] > 0.05, "Positive")
    .when(df_sentiment["sentiment_score"] < -0.05, "Negative")
    .otherwise("Neutral")
)
```

```
df_sentiment.select("reviews", "sentiment_score", "sentiment_label").show(10)
```

↳

reviews	sentiment_score	sentiment_label
---------	-----------------	-----------------

The pizza, steak ...	0.6114	Positive
Great way to get ...	0.6249	Positive
Went there for a ...	0.7392	Positive
Friendly and know...	0.4939	Positive
My personal choic...	0.8308	Positive
This store has be...	-0.6324	Negative
I am disappointed...	0.975	Positive
AUAF oversee a va...	0.9042	Positive
Great neighborhood...	0.6249	Positive
This is place is ...	0.8807	Positive

```
+-----+
only showing top 10 rows
```

```
from pyspark.sql.functions import approx_count_distinct
```

```
df_sentiment_counts = df_sentiment.groupBy("sentiment_label").agg(approx_count_distinct("gmap_id").alias("approx_count"))
```

```
from pyspark.sql.functions import approx_count_distinct
```

```
df_sentiment.groupBy("sentiment_label").agg(approx_count_distinct("gmap_id").alias("approx_count")).show()
```

```
↩ +-----+
|sentiment_label|approx_count|
+-----+
|      Positive|      2647752|
|      Neutral|       91804|
|      Negative|      227790|
+-----+
```

```
from pyspark.sql.functions import col, round
```

```
df_percentages = df_sentiment_counts.withColumn(
    "percentage", round((col("approx_count") / 2884722) * 100, 2)
)
```

✓ TextBlob

```
from pyspark.sql.functions import col, when, regexp_replace
from textblob import TextBlob
from pyspark.sql.functions import udf
from pyspark.sql.types import StructType, StructField, DoubleType, StringType
```

```
def get_textblob_sentiment(text):
    if text is not None and isinstance(text, str):
        blob = TextBlob(text)
        return float(blob.sentiment.polarity), float(blob.sentiment.subjectivity)
    else:
        return 0.0, 0.0
```

```
schema = StructType([
    StructField("polarity", DoubleType(), True),
    StructField("subjectivity", DoubleType(), True)
])
```

```
sentiment_udf = udf(get_textblob_sentiment, schema)
```

```
df_reviews = df_reviews.withColumn("sentiment", sentiment_udf("reviews"))
df_reviews = df_reviews.withColumn("polarity", col("sentiment.polarity"))
df_reviews = df_reviews.withColumn("subjectivity", col("sentiment.subjectivity"))
```

```
from pyspark.sql.functions import when
```

```
# Putting all the reviews into 3 categories based on their polarity
df_reviews = df_reviews.withColumn(
    "sentiment_label",
    when(col("polarity") <= -0.2, "Negative").
    when(col("polarity") <= 0.2, "Neutral").
    otherwise("Positive")
)
```

```
df_reviews.select('reviews', 'polarity', 'sentiment_label').show(10, truncate=False)
```

```

↩ +-----+
  |reviews|
  +-----+
  |The pizza, steak street tacos and the portobello truffle fries and were all delicious!
  |Great way to get around town, and very affordable.
  |Went there for a volleyball tournament. The facilities are nice. They have a lot of bleacher seating. They offered guest wifi but I di
  |Friendly and knowledgeable. I was able to get a copy of my pets shots.
  |My personal choice for somewhat upscale dining in Woodstock. The burgers are delicious, but my go to item is the Lobster Cobb Salad. Yc
  |This store has been featured on the Instagram account of "overpriced bourbon". Tells you all you need to know. Along with the miserable
  |I am disappointed to see negative reviews for this establishment. I am very frequent here and always have a pleasant experience. The be
  |AUAF oversee a variety programs. I work as a Home Care Aide through them. As an employee of AUAF Home Care program I was treated with
  |Great neighborhood bar
  |This is place is solid. The prices are reasonable and the portions large. However the food is not stellar. Some things are pretty good,
  +-----+
only showing top 10 rows

```

```
sentiment_label_count = df_reviews.groupby('sentiment_label').count()
```

```

from pyspark.sql.functions import col, round
df_sentiment_all = df_reviews.withColumn("sentiment_score", vader_udf("reviews"))
df_sentiment_all = df_sentiment_all.withColumn("polarity", col("sentiment.polarity"))
df_sentiment_all.columns

```

```

↩ ['gmap_id',
  'customer_name',
  'rating',
  'reviews',
  'time',
  'avg_rating',
  'category',
  'latitude',
  'longitude',
  'business_name',
  'num_of_reviews',
  'state',
  'standard_category',
  'Monday',
  'Tuesday',
  'Wednesday',
  'Thursday',
  'Friday',
  'Saturday',
  'Sunday',
  'timestamp',
  'week',
  'month',
  'year',
  'time_seconds',
  'sentiment',
  'polarity',
  'subjectivity',
  'sentiment_label',
  'sentiment_score']

```

```

from pyspark.sql.functions import when, length, col
from pyspark.sql import functions as F

```

```

df_labeled = df_sentiment_all.withColumn(
    "vader_label",
    when(col("sentiment_score") >= 0.05, "Positive")
    .when(col("sentiment_score") <= -0.05, "Negative")
    .otherwise("Neutral")
)
df_labeled = df_labeled.withColumn(
    "textblob_label",
    when(col("polarity") >= 0.05, "Positive")
    .when(col("polarity") <= -0.05, "Negative")
    .otherwise("Neutral")
)
df_labeled = df_labeled.withColumn("review_length", length(col("reviews")))

df_labeled = df_labeled.withColumn(
    "length_bucket",
    when(col("review_length") < 100, "Short")
    .when(col("review_length") <= 300, "Medium")

```

```

        .otherwise("Long")
    )

top_categories = [row['standard_category'] for row in df_labeled.groupBy("standard_category")
                  .count().orderBy(F.desc("count")).limit(5).collect()]

df_filtered = df_labeled.filter(col("standard_category").isin(top_categories))

samples = []
for category in top_categories:
    for label in ["Positive", "Negative", "Neutral"]:
        for length_group in ["Short", "Medium", "Long"]:
            subset = (
                df_filtered.filter(
                    (col("standard_category") == category) &
                    (col("vader_label") == label) &
                    (col("length_bucket") == length_group)
                ).orderBy(F.rand()).limit(1)
            )
            samples.append(subset)

df_sample_30 = samples[0]
for i in range(1, len(samples)):
    df_sample_30 = df_sample_30.union(samples[i])

df_final = df_sample_30.select(
    "reviews", "standard_category", "review_length",
    "vader_label", "sentiment", "textblob_label", "polarity"
)

df_final.show(30, truncate=False)

df_final_30 = df_final.limit(30)

from transformers import pipeline

# Load zero-shot classification model
zero_shot_classifier = pipeline("zero-shot-classification", model="facebook/bart-large-mnli")

from pyspark.sql.functions import when, length, col

df_bucketed = df_reviews.withColumn("review_length", length(col("reviews")))

df_bucketed = df_bucketed.withColumn(
    "length_bucket",
    when(col("review_length") < 100, "Short")
    .when(col("review_length") <= 300, "Medium")
    .otherwise("Long")
)

from pyspark.sql import functions as F

# Group count
group_counts = df_bucketed.groupBy("standard_category", "length_bucket").count()

# Total rows
total_count = df_bucketed.count()

# Target sample size
sample_target = 10000

# Compute fraction per group
group_fractions = group_counts.withColumn(
    "fraction", (F.col("count") / total_count) * sample_target
).withColumn(
    "sample_size", F.round("fraction").cast("int")
)

sampled_dfs = []

for row in group_fractions.collect():
    cat = row['standard_category']

```

```

bucket = row['length_bucket']
n = row['sample_size']

if n > 0:
    subset = (
        df_bucketed.filter(
            (col("standard_category") == cat) &
            (col("length_bucket") == bucket)
        )
        .orderBy(F.rand())
        .limit(n)
    )
    sampled_dfs.append(subset)

df_sample_10k = sampled_dfs[0]
for sdf in sampled_dfs[1:]:
    df_sample_10k = df_sample_10k.union(sdf)

output_path = "gs://final_dataset_dat490/sample_reviews_stratified_10k.parquet"

df_sample_10k.write.mode("overwrite").parquet(output_path)

import pandas as pd
df_sample_pandas = pd.read_parquet(output_path)
df_sample_pandas = df_sample_pandas.sort_values(["gmap_id", "timestamp"]).reset_index(drop=True)
df_sample_pandas.columns

Index(['gmap_id', 'customer_name', 'rating', 'reviews', 'time', 'avg_rating',
      'category', 'latitude', 'longitude', 'business_name', 'num_of_reviews',
      'state', 'standard_category', 'Monday', 'Tuesday', 'Wednesday',
      'Thursday', 'Friday', 'Saturday', 'Sunday', 'timestamp', 'week',
      'month', 'year', 'time_seconds', 'review_length', 'length_bucket'],
      dtype='object')

```

```
df_sample_pandas[['reviews']]
```

↗

	reviews
0	Nice clean place and very friendly staff. Only...
1	This company helped us buy our land 15 years a...
2	Great place for kids and families to do sports...
3	He was able to fit me in at the last minute. A...
4	I have had good experiences here for the past ...
...	...
9996	Store was disorganized and sales representativ...
9997	Treated as a number, not a patient. Rushed by ...
9998	This is such a great place to grab dinner. The...
9999	Dividers and spots clearly marked on floor, pl...
10000	Travelers, this is the cape cod thrift store e...

10001 rows x 1 columns

◀ ▶

```

# pip install pandas transformers openpyxl

# from google.colab import drive
# drive.mount('/content/drive')

# import pandas as pd
# df_reviews_30 = pd.read_excel('/content/drive/MyDrive/Reviews std DAT490.xlsx')
# df_reviews_30.head()

# from transformers import pipeline

# classifier = pipeline("zero-shot-classification", model="facebook/bart-large-mnli")
# labels = ["positive", "neutral", "negative"]

```

```
# score = []
# for review in df_reviews_30['Reviews']:
#     prediction = classifier(review, candidate_labels=labels)
#     top_label = prediction["labels"][0]
#     score.append(top_label)

# df_reviews_30['Sentiment'] = score
# df_reviews_30

# df_reviews_30.to_excel('/content/drive/MyDrive/Reviews std DAT490 labelled.xlsx')
```

```
!pip install vaderSentiment
```

```
Collecting vaderSentiment
  Downloading vaderSentiment-3.3.2-py2.py3-none-any.whl.metadata (572 bytes)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from vaderSentiment) (2.32.3)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->vaderSentiment) (3.4.
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests->vaderSentiment) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests->vaderSentiment) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests->vaderSentiment) (2025.1.31)
Downloading vaderSentiment-3.3.2-py2.py3-none-any.whl (125 kB)
126.0/126.0 kB 4.0 MB/s eta 0:00:00
Installing collected packages: vaderSentiment
Successfully installed vaderSentiment-3.3.2
```

```
from textblob import TextBlob
from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer
from tqdm import tqdm
```

```
tqdm.pandas() # Progress bar for apply
```

```
analyzer = SentimentIntensityAnalyzer()
```

```
def get_textblob_polarity(text):
    try:
        return TextBlob(text).sentiment.polarity
    except:
        return None
```

```
def get_vader_compound(text):
    try:
        return analyzer.polarity_scores(text)["compound"]
    except:
        return None
```

```
# Assuming df["reviews"] contains your text
df_sample_pandas["vader_polarity"] = df_sample_pandas["reviews"].progress_apply(get_textblob_polarity)
df_sample_pandas["textblob_sentiment_score"] = df_sample_pandas["reviews"].progress_apply(get_vader_compound)
```

```
100% |██████████| 10001/10001 [00:03<00:00, 2803.51it/s]
100% |██████████| 10001/10001 [00:03<00:00, 2917.04it/s]
```

```
def label_textblob(p):
    if p >= 0.05:
        return "Positive"
    elif p <= -0.05:
        return "Negative"
    else:
        return "Neutral"
```

```
def label_vader(s):
    if s >= 0.05:
        return "Positive"
    elif s <= -0.05:
        return "Negative"
    else:
        return "Neutral"
```

```
# Updated column names
df_sample_pandas["textblob_label"] = df_sample_pandas["vader_polarity"].apply(label_textblob)
```

```
df_sample_pandas["vader_label"] = df_sample_pandas["textblob_sentiment_score"].apply(label_vader)
df_sample_pandas.head()
```




	gmap_id	customer_name	rating	reviews	time	avg_rating	category	latitude	longitude
0	0x0:0xde4ab363e58baf8	Glen Sikorski	5	Nice clean place and very friendly staff. Only...	1619963856471	4.2	Restaurant, Bar, Caterer, Event venue	45.120739	-91.448845
1	0x145e95d513a77c99:0x7aad3c9a54c17e9f	stephen travers	5	This company helped us buy our land 15 years a...	1621983091409	3.4	Property management company, Real estate agency	29.730017	-99.075182
2	0x14e037302ebfe6bd:0x483c80e39ebb0ab7	Shawn Bebej	5	Great place for kids and families to do sports...	1578142093772	4.9	Non-profit organization	34.233393	-118.590456



```
from transformers import pipeline
from tqdm import tqdm

# Enable progress bar
tqdm.pandas()

# Load zero-shot classification pipeline with BART
zero_shot_classifier = pipeline("zero-shot-classification", model="facebook/bart-large-mnli")
```



config.json: 100%

1.15k/1.15k [00:00<00:00, 74.2kB/s]

Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download. For better performance, please install the 'hf_xet' package. See [https://huggingface.co/docs/xet/en/index](#) for more details.

model.safetensors: 100%

1.63G/1.63G [00:10<00:00, 244MB/s]

tokenizer_config.json: 100%

26.0/26.0 [00:00<00:00, 2.94kB/s]

vocab.json: 100%


899k/899k [00:00<00:00, 21.7MB/s]

merges.txt: 100%

456k/456k [00:01<00:00, 335kB/s]

tokenizer.json: 100%


1.36M/1.36M [00:00<00:00, 30.7MB/s]



```
candidate_labels = ["positive", "neutral", "negative"]

def classify_bart(text):
    try:
        result = zero_shot_classifier(text, candidate_labels)
        return result["labels"][0].capitalize() # Most likely label
    except:
        return None # Handle any errors gracefully

df_sample_pandas["bart_label"] = df_sample_pandas["reviews"].progress_apply(classify_bart)
```



0%|██████████| 10/10001 [00:02<22:20, 7.45it/s]

100%|██████████| 10001/10001 [14:47<00:00, 11.27it/s]

You seem to be using the pipelines sequentially on GPU. In order to maximize efficiency, please use DataParallel.

df_sample_pandas.head()

		gmap_id	customer_name	rating	reviews	time	avg_rating	category	latitude	longitude
0		0x0:0xde4ab363e58baf8	Glen Sikorski	5	Nice clean place and very friendly staff. Only...	1619963856471	4.2	Restaurant, Bar, Caterer, Event venue	45.120739	-91.448845
1	0x145e95d513a77c99:0x7aad3c9a54c17e9f		stephen travers	5	This company helped us buy our land 15 years a...	1621983091409	3.4	Property management company, Real estate agency	29.730017	-99.075182
2	0x14e037302ebfe6bd:0x483c80e39ebb0ab7		Shawn Bebej	5	Great place for kids and families to do sports...	1578142093772	4.9	Non-profit organization	34.233393	-118.590456
3	0x14e3db41cf753ebd:0x1d6536e7c20051ef		L C	5	He was able to fit me in at the	1616604264474	4.8	Window tinting	33.546954	-112.202525

```
import pandas as pd

# Count label frequencies per method
label_counts = pd.DataFrame({
    "TextBlob": df_sample_pandas["textblob_label"].value_counts(),
    "VADER": df_sample_pandas["vader_label"].value_counts(),
    "BART": df_sample_pandas["bart_label"].value_counts()
}).fillna(0).astype(int)

# Reorder rows for consistency
label_counts = label_counts.reindex(["Positive", "Neutral", "Negative"])

# Melt into long format
df_melted = label_counts.T.reset_index().melt(
    id_vars="index", var_name="Sentiment", value_name="Count"
)
df_melted.rename(columns={"index": "Method"}, inplace=True)

import seaborn as sns
import matplotlib.pyplot as plt

# Set seaborn style
sns.set(style="whitegrid")

# Define consistent color palette
palette = {
    "Positive": "#4CAF50",
    "Neutral": "#FFC107",
    "Negative": "#F44336"
}

# Create grouped bar plot
plt.figure(figsize=(10, 6))
sns.barplot(
    data=df_melted,
    x="Method", y="Count", hue="Sentiment",
```

```
palette=palette
)

# Customize the chart
plt.title("Sentiment Label Distribution (10,000 Reviews)", fontsize=14)
plt.xlabel("Sentiment Analysis Method")
plt.ylabel("Number of Reviews")
plt.legend(title="Sentiment")
plt.tight_layout()
plt.show()
```

