

Methodology

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
```

```
In [ ]: from google.colab import auth
auth.authenticate_user()
```

```
In [ ]: from google.cloud import storage

project_id = "sharp-matter-449521-u2"
!gcloud config set project {project_id}
```

Updated property [core/project].

```
In [ ]: !wget -P /usr/lib/spark/jars/ https://storage.googleapis.com/hadoop-lib/gcs/gcs-connector-hadoop3-latest.jar

--2025-04-09 22:20:49-- https://storage.googleapis.com/hadoop-lib/gcs/gcs-connector-hadoop3-latest.jar
Resolving storage.googleapis.com (storage.googleapis.com)... 172.253.117.207, 142.250.99.207, 142.250.107.207, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|172.253.117.207|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 40713341 (39M) [application/java-archive]
Saving to: '/usr/lib/spark/jars/gcs-connector-hadoop3-latest.jar'

gcs-connector-hadoo 100%[=====>] 38.83M 207MB/s in 0.2s

2025-04-09 22:20:49 (207 MB/s) - '/usr/lib/spark/jars/gcs-connector-hadoop3-latest.jar' saved [40713341/40713341]
```

```
In [ ]: import warnings
warnings.filterwarnings('ignore')
```

```
In [ ]: from pyspark.sql import SparkSession
```

```
In [ ]: spark = SparkSession.builder \
    .appName("BigDataProcessing") \
    .config("spark.jars", "/usr/lib/spark/jars/gcs-connector-hadoop3-latest.jar") \
    .config("spark.hadoop.fs.gs.impl", "com.google.cloud.hadoop.fs.gcs.GoogleHadoopFileSystem") \
    .config("spark.hadoop.fs.gs.auth.service.account.enable", "true") \
    .getOrCreate()
```

```
In [ ]: spark
```

Out []: **SparkSession - in-memory**

SparkContext

Spark UI

Version	v3.5.5
Master	local[*]
AppName	BigDataProcessing

```
In [ ]: df_reviews = spark.read.parquet('gs://final_dataset_dat490/dat490_final_datase'
```

```
In [ ]: df_reviews.columns
```

```
Out [ ]: ['gmap_id',  
         'customer_name',  
         'rating',  
         'reviews',  
         'time',  
         'avg_rating',  
         'category',  
         'latitude',  
         'longitude',  
         'business_name',  
         'num_of_reviews',  
         'state',  
         'standard_category',  
         'Monday',  
         'Tuesday',  
         'Wednesday',  
         'Thursday',  
         'Friday',  
         'Saturday',  
         'Sunday',  
         'timestamp',  
         'week',  
         'month',  
         'year',  
         'time_seconds']
```

VADER

```
In [ ]: from pyspark.sql.functions import udf  
from pyspark.sql.types import FloatType  
from nltk.sentiment.vader import SentimentIntensityAnalyzer  
import nltk  
nltk.download("vader_lexicon")  
  
# Initializing VADER  
sia = SentimentIntensityAnalyzer()  
  
def vader_sentiment(text):  
    if text:  
        return float(sia.polarity_scores(text)["compound"])  
    else:  
        return 0.0
```

```
vader_udf = udf(vader_sentiment, FloatType())

df_sentiment = df_reviews.withColumn("sentiment_score", vader_udf("reviews"))
```

[nltk_data] Downloading package vader_lexicon to /root/nltk_data...

```
In [ ]: from pyspark.sql.functions import when

df_sentiment = df_sentiment.withColumn(
    "sentiment_label",
    when(df_sentiment["sentiment_score"] > 0.2, "Positive")
    .when(df_sentiment["sentiment_score"] < -0.2, "Negative")
    .otherwise("Neutral")
)
```

```
In [ ]: from pyspark.sql.functions import approx_count_distinct

df_sentiment.groupBy("sentiment_label").agg(approx_count_distinct("gmap_id").a
```

sentiment_label	approx_count
Positive	2598785
Neutral	149531
Negative	197929

```
In [ ]: from pyspark.sql.functions import when

df_sentiment = df_sentiment.withColumn(
    "sentiment_label",
    when(df_sentiment["sentiment_score"] > 0.05, "Positive")
    .when(df_sentiment["sentiment_score"] < -0.05, "Negative")
    .otherwise("Neutral")
)
```

```
In [ ]: df_sentiment.select("reviews", "sentiment_score", "sentiment_label").show(10)
```

reviews	sentiment_score	sentiment_label
The pizza, steak ...	0.6114	Positive
Great way to get ...	0.6249	Positive
Went there for a ...	0.7392	Positive
Friendly and know...	0.4939	Positive
My personal choic...	0.8308	Positive
This store has be...	-0.6324	Negative
I am disappointed...	0.975	Positive
AUAF oversee a va...	0.9042	Positive
Great neighborhoo...	0.6249	Positive
This is place is ...	0.8807	Positive

only showing top 10 rows

```
In [ ]: from pyspark.sql.functions import approx_count_distinct

df_sentiment_counts = df_sentiment.groupBy("sentiment_label").agg(approx_count
```

```
In [ ]: from pyspark.sql.functions import approx_count_distinct

df_sentiment.groupBy("sentiment_label").agg(approx_count_distinct("gmap_id").a

+-----+-----+
|sentiment_label|approx_count|
+-----+-----+
|      Positive|      2647752|
|      Neutral|       91804|
|      Negative|      227790|
+-----+-----+
```

```
In [ ]: from pyspark.sql.functions import col, round

df_percentages = df_sentiment_counts.withColumn(
    "percentage", round((col("approx_count") / 2884722) * 100, 2)
)
```

TextBlob

```
In [ ]: from pyspark.sql.functions import col, when, regexp_replace
from textblob import TextBlob
from pyspark.sql.functions import udf
from pyspark.sql.types import StructType, StructField, DoubleType, StringType

def get_textblob_sentiment(text):
    if text is not None and isinstance(text, str):
        blob = TextBlob(text)
        return float(blob.sentiment.polarity), float(blob.sentiment.subjectivity)
    else:
        return 0.0, 0.0

schema = StructType([
    StructField("polarity", DoubleType(), True),
    StructField("subjectivity", DoubleType(), True)
])

sentiment_udf = udf(get_textblob_sentiment, schema)

df_reviews = df_reviews.withColumn("sentiment", sentiment_udf("reviews"))
df_reviews = df_reviews.withColumn("polarity", col("sentiment.polarity"))
df_reviews = df_reviews.withColumn("subjectivity", col("sentiment.subjectivity"))
```

```
In [ ]: from pyspark.sql.functions import when

# Putting all the reviews into 3 categories based on their polarity
df_reviews = df_reviews.withColumn(
    "sentiment_label",
    when(col("polarity") <= -0.2, "Negative").
    when(col("polarity") <= 0.2, "Neutral").
    otherwise("Positive")
)
```

```
In [ ]: df_reviews.select('reviews', 'polarity', 'sentiment_label').show(10, truncate=1)
```

5/15

an employee of AUAF Home Care program I was treated with respect and dignity as a Home Aide.

0.0	Neutral
-----	---------

|Great neighborhood bar

0.8	Positive
-----	----------

This place is solid. The prices are reasonable and the portions large. However the food is not stellar. Some things are pretty good, others leave something to be desired in terms of flavorful sauces, etc. I had a frustrating customer service experience with delivery once and the owner was too frustrated/busy to resolve the issue so GrubHub refunded us. It's a good local spot.

0.171031746031746	Neutral
-------------------	---------

This image shows a blank sheet of primary-ruled paper. It features ten horizontal dashed lines spaced evenly down the page. On the left side, there are two solid vertical lines forming a margin; the first is near the left edge, and the second is further right. On the right side, there is one solid vertical line. Small black tick marks are located at the intersections of the horizontal lines with the leftmost and rightmost vertical lines. The central area between the vertical lines is intended for writing.

only showing top 10 rows

```
In [ ]: sentiment_label_count = df_reviews.groupby('sentiment_label').count()
```

```
In [ ]: from pyspark.sql.functions import col, round
df_sentiment_all = df_reviews.withColumn("sentiment_score", vader_udf("reviews"))
df_sentiment_all = df_sentiment_all.withColumn("polarity", col("sentiment.polarity"))
df_sentiment_all.columns
```

```
Out[ ]: ['gmap_id',
        'customer_name',
        'rating',
        'reviews',
        'time',
        'avg_rating',
        'category',
        'latitude',
        'longitude',
        'business_name',
        'num_of_reviews',
        'state',
        'standard_category',
        'Monday',
        'Tuesday',
        'Wednesday',
        'Thursday',
        'Friday',
        'Saturday',
        'Sunday',
        'timestamp',
        'week',
        'month',
        'year',
        'time_seconds',
        'sentiment',
        'polarity',
        'subjectivity',
        'sentiment_label',
        'sentiment_score']
```

```
In [ ]: from pyspark.sql.functions import when, length, col
        from pyspark.sql import functions as F

        df_labeled = df_sentiment_all.withColumn(
            "vader_label",
            when(col("sentiment_score") >= 0.05, "Positive")
            .when(col("sentiment_score") <= -0.05, "Negative")
            .otherwise("Neutral")
        )
        df_labeled = df_labeled.withColumn(
            "textblob_label",
            when(col("polarity") >= 0.05, "Positive")
            .when(col("polarity") <= -0.05, "Negative")
            .otherwise("Neutral")
        )
        df_labeled = df_labeled.withColumn("review_length", length(col("reviews")))

        df_labeled = df_labeled.withColumn(
            "length_bucket",
            when(col("review_length") < 100, "Short")
            .when(col("review_length") <= 300, "Medium")
            .otherwise("Long")
        )

        top_categories = [row['standard_category'] for row in df_labeled.groupBy("standard_category")
                           .count().orderBy(F.desc("count")).limit(5).collect()]

        df_filtered = df_labeled.filter(col("standard_category").isin(top_categories))
```

```

samples = []
for category in top_categories:
    for label in ["Positive", "Negative", "Neutral"]:
        for length_group in ["Short", "Medium", "Long"]:
            subset = (
                df_filtered.filter(
                    (col("standard_category") == category) &
                    (col("vader_label") == label) &
                    (col("length_bucket") == length_group)
                ).orderBy(F.rand()).limit(1)
            )
            samples.append(subset)

df_sample_30 = samples[0]
for i in range(1, len(samples)):
    df_sample_30 = df_sample_30.union(samples[i])

df_final = df_sample_30.select(
    "reviews", "standard_category", "review_length",
    "vader_label", "sentiment", "textblob_label", "polarity"
)

df_final.show(30, truncate=False)

```

```
In [ ]: df_final_30 = df_final.limit(30)
```

```
In [ ]: from transformers import pipeline

# Load zero-shot classification model
zero_shot_classifier = pipeline("zero-shot-classification", model="facebook/ba
```

```
In [ ]: from pyspark.sql.functions import when, length, col

df_bucketed = df_reviews.withColumn("review_length", length(col("reviews")))

df_bucketed = df_bucketed.withColumn(
    "length_bucket",
    when(col("review_length") < 100, "Short")
    .when(col("review_length") <= 300, "Medium")
    .otherwise("Long")
)

```

```
In [ ]: from pyspark.sql import functions as F

# Group count
group_counts = df_bucketed.groupBy("standard_category", "length_bucket").count

# Total rows
total_count = df_bucketed.count()

# Target sample size
sample_target = 10000

# Compute fraction per group
group_fractions = group_counts.withColumn(
    "fraction", (F.col("count") / total_count) * sample_target
).withColumn(

```



```
"sample_size", F.round("fraction").cast("int")
)
```

```
In [ ]: sampled_dfs = []

for row in group_fractions.collect():
    cat = row['standard_category']
    bucket = row['length_bucket']
    n = row['sample_size']

    if n > 0:
        subset = (
            df_bucketed.filter(
                (col("standard_category") == cat) &
                (col("length_bucket") == bucket)
            )
            .orderBy(F.rand())
            .limit(n)
        )
        sampled_dfs.append(subset)
```

```
In [ ]: df_sample_10k = sampled_dfs[0]
for sdf in sampled_dfs[1:]:
    df_sample_10k = df_sample_10k.union(sdf)
```

```
In [ ]: output_path = "gs://final_dataset_dat490/sample_reviews_stratified_10k.parquet"
```

```
In [ ]: df_sample_10k.write.mode("overwrite").parquet(output_path)
```

```
In [ ]: import pandas as pd
df_sample_pandas = pd.read_parquet(output_path)
df_sample_pandas = df_sample_pandas.sort_values(["gmap_id", "timestamp"]).reset_index()
df_sample_pandas.columns
```

```
Out[ ]: Index(['gmap_id', 'customer_name', 'rating', 'reviews', 'time', 'avg_rating',
        'category', 'latitude', 'longitude', 'business_name', 'num_of_reviews',
        'state', 'standard_category', 'Monday', 'Tuesday', 'Wednesday',
        'Thursday', 'Friday', 'Saturday', 'Sunday', 'timestamp', 'week',
        'month', 'year', 'time_seconds', 'review_length', 'length_bucket'],
        dtype='object')
```

```
In [ ]: df_sample_pandas[['reviews']]
```

Out []:

reviews

0	Nice clean place and very friendly staff. Only...
1	This company helped us buy our land 15 years a...
2	Great place for kids and families to do sports...
3	He was able to fit me in at the last minute. A...
4	I have had good experiences here for the past ...
...	...
9996	Store was disorganized and sales representativ...
9997	Treated as a number, not a patient. Rushed by ...
9998	This is such a great place to grab dinner. The...
9999	Dividers and spots clearly marked on floor, pl...
10000	Travelers, this is the cape cod thrift store e...

10001 rows × 1 columns

```
In [ ]: # pip install pandas transformers openpyxl

# from google.colab import drive
# drive.mount('/content/drive')

# import pandas as pd
# df_reviews_30 = pd.read_excel('/content/drive/MyDrive/Reviews std DAT490.xls')
# df_reviews_30.head()

# from transformers import pipeline

# classifier = pipeline("zero-shot-classification", model="facebook/bart-large-
# labels = ["positive", "neutral", "negative"]

# score = []
# for review in df_reviews_30['Reviews']:
#     prediction = classifier(review, candidate_labels=labels)
#     top_label = prediction["labels"][0]
#     score.append(top_label)

# df_reviews_30['Sentiment'] = score
# df_reviews_30

# df_reviews_30.to_excel('/content/drive/MyDrive/Reviews std DAT490 labelled.xls')
```

```
In [ ]: !pip install vaderSentiment
```

Collecting vaderSentiment

Downloading vaderSentiment-3.3.2-py2.py3-none-any.whl.metadata (572 bytes)
 Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from vaderSentiment) (2.32.3)
 Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->vaderSentiment) (3.4.1)
 Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests->vaderSentiment) (3.10)
 Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests->vaderSentiment) (2.3.0)
 Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests->vaderSentiment) (2025.1.31)
 Downloading vaderSentiment-3.3.2-py2.py3-none-any.whl (125 kB)

126.0/126.0 kB 4.0 MB/s eta 0:00:00

Installing collected packages: vaderSentiment
 Successfully installed vaderSentiment-3.3.2

```
In [ ]: from textblob import TextBlob
        from vaderSentiment.vaderSentiment import SentimentIntensityAnalyzer
        from tqdm import tqdm

        tqdm.pandas() # Progress bar for apply
```

```
In [ ]: analyzer = SentimentIntensityAnalyzer()

        def get_textblob_polarity(text):
            try:
                return TextBlob(text).sentiment.polarity
            except:
                return None

        def get_vader_compound(text):
            try:
                return analyzer.polarity_scores(text)["compound"]
            except:
                return None
```

```
In [ ]: # Assuming df["reviews"] contains your text
        df_sample_pandas["vader_polarity"] = df_sample_pandas["reviews"].progress_apply(
        df_sample_pandas["textblob_sentiment_score"] = df_sample_pandas["reviews"].progress_apply(
```

```
100%|██████████| 10001/10001 [00:03<00:00, 2803.51it/s]
100%|██████████| 10001/10001 [00:03<00:00, 2917.04it/s]
```

```
In [ ]: def label_textblob(p):
        if p >= 0.05:
            return "Positive"
        elif p <= -0.05:
            return "Negative"
        else:
            return "Neutral"

        def label_vader(s):
            if s >= 0.05:
                return "Positive"
            elif s <= -0.05:
                return "Negative"
            else:
```

```
return "Neutral"

# Updated column names
df_sample_pandas["textblob_label"] = df_sample_pandas["vader_polarity"].apply(lambda x: "Positive" if x > 0 else "Negative" if x < 0 else "Neutral")
df_sample_pandas["vader_label"] = df_sample_pandas["textblob_sentiment_score"].apply(lambda x: "Positive" if x > 0 else "Negative" if x < 0 else "Neutral")
df_sample_pandas.head()
```

Out []:

	gmap_id	customer_name	rating	reviews	tim
0	0x0:0xde4ab363e58baf8	Glen Sikorski	5	Nice clean place and very friendly staff. Only...	161996385647
1	0x145e95d513a77c99:0x7aad3c9a54c17e9f	stephen travers	5	This company helped us buy our land 15 years a...	162198309140
2	0x14e037302ebfe6bd:0x483c80e39ebb0ab7	Shawn Bebej	5	Great place for kids and families to do sports...	157814209377
3	0x14e3db41cf753ebd:0x1d6536e7c20051ef	L C	5	He was able to fit me in at the last minute. A...	161660426447
4	0x1520f8e750be33d7:0xc5e501b57143e755	Lorelei Flaherty	4	I have had good experiences here for the past ...	147707233430

5 rows x 31 columns

In []:

```
from transformers import pipeline
from tqdm import tqdm

# Enable progress bar
tqdm.pandas()

# Load zero-shot classification pipeline with BART
zero_shot_classifier = pipeline("zero-shot-classification", model="facebook/bart-large-mnli")

config.json: 0%|          | 0.00/1.15k [00:00<?, ?B/s]
Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download. For better performance, install the package with: `pip install huggingface_hub[hf_xet]` or `pip install hf_xet`
WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download. For better performance, install the package with: `pip install huggingface_hub[hf_xet]` or `pip install hf_xet`
model.safetensors: 0%|          | 0.00/1.63G [00:00<?, ?B/s]
tokenizer_config.json: 0%|          | 0.00/26.0 [00:00<?, ?B/s]
vocab.json: 0%|          | 0.00/899k [00:00<?, ?B/s]
```

merges.txt: 0%| | 0.00/456k [00:00<?, ?B/s]

tokenizer.json: 0%| | 0.00/1.36M [00:00<?, ?B/s]

Device set to use cuda:0

In []:

```
candidate_labels = ["positive", "neutral", "negative"]

def classify_bart(text):
    try:
        result = zero_shot_classifier(text, candidate_labels)
        return result["labels"][0].capitalize() # Most likely label
    except:
        return None # Handle any errors gracefully
```

In []:

```
df_sample_pandas["bart_label"] = df_sample_pandas["reviews"].progress_apply(classify_bart)

0%| | 10/10001 [00:02<22:20, 7.45it/s]You seem to be using the pipelines sequentially on GPU. In order to maximize efficiency please use a dataset
100%|██████████| 10001/10001 [14:47<00:00, 11.27it/s]
```

In []:

```
df_sample_pandas.head()
```

Out []:

	gmap_id	customer_name	rating	reviews	time
0	0x0:0xde4ab363e58baf8	Glen Sikorski	5	Nice clean place and very friendly staff. Only...	161996385647
1	0x145e95d513a77c99:0x7aad3c9a54c17e9f	stephen travers	5	This company helped us buy our land 15 years a...	162198309140
2	0x14e037302ebfe6bd:0x483c80e39ebb0ab7	Shawn Bebej	5	Great place for kids and families to do sports...	157814209377
3	0x14e3db41cf753ebd:0x1d6536e7c20051ef	L C	5	He was able to fit me in at the last minute. A...	161660426447
4	0x1520f8e750be33d7:0xc5e501b57143e755	Lorelei Flaherty	4	I have had good experiences here for the past ...	147707233430

5 rows × 32 columns

In []:

```
import pandas as pd

# Count label frequencies per method
label_counts = pd.DataFrame({
    "TextBlob": df_sample_pandas["textblob_label"].value_counts(),
```

```
"VADER": df_sample_pandas["vader_label"].value_counts(),
"BART": df_sample_pandas["bart_label"].value_counts()
}).fillna(0).astype(int)

# Reorder rows for consistency
label_counts = label_counts.reindex(["Positive", "Neutral", "Negative"])
```

```
In [ ]: # Melt into long format
df_melted = label_counts.T.reset_index().melt(
    id_vars="index", var_name="Sentiment", value_name="Count"
)
df_melted.rename(columns={"index": "Method"}, inplace=True)
```

```
In [ ]: import seaborn as sns
import matplotlib.pyplot as plt

# Set seaborn style
sns.set(style="whitegrid")

# Define consistent color palette
palette = {
    "Positive": "#4CAF50",
    "Neutral": "#FFC107",
    "Negative": "#F44336"
}

# Create grouped bar plot
plt.figure(figsize=(10, 6))
sns.barplot(
    data=df_melted,
    x="Method", y="Count", hue="Sentiment",
    palette=palette
)

# Customize the chart
plt.title("Sentiment Label Distribution (10,000 Reviews)", fontsize=14)
plt.xlabel("Sentiment Analysis Method")
plt.ylabel("Number of Reviews")
plt.legend(title="Sentiment")
plt.tight_layout()
plt.show()
```



```
In [ ]: import json
# Change this to the name of your broken notebook
notebook_filename = "your_notebook_name.ipynb"
# Load the notebook
with open(notebook_filename, 'r', encoding='utf-8') as f:
    notebook_data = json.load(f)
# Fix metadata.widgets if missing 'state'
widgets = notebook_data.get('metadata', {}).get('widgets', {})
if 'application/vnd.jupyter.widget-state+json' in widgets:
    widget_meta = widgets['application/vnd.jupyter.widget-state+json']
    if 'state' not in widget_meta:
        widget_meta['state'] = {}
        widget_meta['version_major'] = 2
        widget_meta['version_minor'] = 0
        print(":white_check_mark: 'state' key added to metadata.widgets.")
else:
    print(":information_source: No widget metadata found or already fixed.")
# Save the fixed notebook (overwrites the original!)
with open(notebook_filename, 'w', encoding='utf-8') as f:
    json.dump(notebook_data, f, indent=2)
print(f":white_check_mark: Notebook '{notebook_filename}' fixed.")
```